

Oluwatobi Melvyn Mayungbo

Machine Learning Engineer, Fraud and Risk

melvyn.t@gmail.com | [LinkedIn: linkedin.com/in/oluwatobi-mayungbo-3a567026b](https://www.linkedin.com/in/oluwatobi-mayungbo-3a567026b) | [GitHub: github.com/MelvTheGoat](https://github.com/MelvTheGoat)

SUMMARY

Machine learning engineer with a B.Sc. in Statistics who takes models from problem framing to production. Built sequence fraud models that beat a strong LightGBM baseline and serve at 5.1 ms p99, a credit decisioning system with calibrated probabilities and a fairness audit, and an uplift study that proved campaign impact.

SKILLS

Machine Learning: PyTorch, LightGBM, Sequence Models (GRU, TCN, Transformers), Model Calibration (Isotonic Regression, Temperature Scaling), SHAP, Feature Engineering, Scikit-learn, XGBoost, Deep Learning

Payments and Risk: Transaction Fraud Detection, Credit Risk Decisioning, Payment Reconciliation, Double-Entry Ledgers and Idempotency, Regulatory Compliance (CBN/NDPA), Demand Forecasting, Web3 Fraud Risk Scoring

Engineering and MLOps: Python, SQL, FastAPI, Docker, Redis, DuckDB, PostgreSQL, GCP Cloud Run, AWS ECS Fargate, Railway, Prefect, dbt-core, MLflow, Evidently, CI/CD (GitHub Actions), Git, Pytest, Streamlit, React

Statistics and Causal Inference: Statistical Inference, Experiment Design, Power Analysis, CUPED, Causal Inference (S/T/X-Learners, Causal Forests), Uplift Modeling, Qini and AUUC, Time-Series Backtesting, Logistic Regression, WOE Scorecards

AI and LLM Systems: Retrieval-Augmented Generation (RAG), Hybrid Search (BM25, Dense Embeddings, RRF), LLM Evaluation and Guardrails, Prompt Engineering, Citation Verification, Model Context Protocol (MCP), Hugging Face, ONNX Runtime

PROJECTS

Sequence-Based Transaction Fraud Detection System | [GitHub](#)

Tech: Python, PyTorch, LightGBM, ONNX Runtime, FastAPI, Docker, SHAP

- Built and compared three deep sequence architectures (GRU, Temporal Convolutional Network and a causally masked Transformer encoder) over customer transaction histories, with learned embeddings for high-cardinality merchant (~2.9k vocabulary), MCC and device-type categoricals; all models under 115k parameters and CPU-trainable in under 6 minutes.
- Benchmarked all architectures against a strong LightGBM baseline on engineered velocity and deviation features, reporting PR-AUC and precision at fixed review capacity (top 0.1%, 0.5%, 1%) across 3 seeds, with a per-attack-type breakdown covering card testing, account takeover and merchant compromise; sequence models cut expected cost per transaction 27% under an explicit cost model.
- Handled extreme class imbalance with focal loss and weighted sampling and isolated their failure modes: focal loss degraded probability calibration (ECE 0.0063) while weighted sampling cost 0.042 PR-AUC. Corrected calibration with temperature scaling and isotonic regression, reducing ECE to 0.0005.
- Deployed real-time scoring by exporting to ONNX Runtime (2.9× faster than eager PyTorch at batch size 1) and serving via FastAPI, reaching 5.1 ms p99 end-to-end latency against a 50 ms budget, with append-only per-decision audit logging and attention-based attribution of the prior transactions driving each flag.

Credit Risk Decisioning and Fairness Audit System | [GitHub](#) | [Live App](#)

Tech: Python, LightGBM, scikit-learn, SHAP, WOE Scorecards, FastAPI, Docker

- Developed a full credit decisioning system producing calibrated default probabilities, cost-optimal approve and decline cutoffs, adverse action reason codes and an append-only audit trail that lets any individual decision be reconstructed months later.
- Benchmarked gradient boosting against a traditional WOE-binned logistic scorecard, prioritising calibration quality (Brier score, reliability diagrams, expected calibration error) over ranking metrics, because expected loss needs true probabilities rather than rankings.
- Audited fairness across sex, age, education and marital status using demographic parity and equal opportunity differences plus within-group calibration, and presented the accuracy and fairness tradeoff curve as a policy decision rather than silently optimising it away.

Uplift Modelling for Campaign Targeting | [GitHub](#)

Tech: Python, scikit-learn, LightGBM, EconML, pytest, GitHub Actions

- Ran a causal inference study on the Hillstrom randomised e-mail trial (64,000 customers) to decide which customers to contact under a fixed budget, delivered as a two-page decision memo for a marketing director.
- Showed the campaign works (1.25% conversion against 0.57%) but that a non-randomised list overstates the effect by 7.5×, and that a response model mails five times as many customers the campaign harms as an uplift model does.

EXPERIENCE

Machine Learning Instructor, SQI College of ICT

March 2026 to Present

- Teach machine learning and deep learning to student cohorts, covering statistical foundations, classical ML algorithms, neural architectures, sequence models and production evaluation metrics.
- Instruct students on data analysis, feature engineering and model building with PyTorch, Scikit-Learn, Pandas and NumPy.

Internal Auditor, NISER

May 2024 to July 2024

- Analysed financial and operational records to identify inconsistencies and irregularities, performing manual anomaly detection across complex transactional datasets.

EDUCATION

B.Sc. Statistics, University of Ibadan

January 2021 to March 2025

Professional Diploma in Artificial Intelligence, SQI College of ICT

July 2025 to Present

CERTIFICATIONS

- Google Data Analytics Professional Certificate
- Google Advanced Data Analytics Professional Certificate
- Claude Code In Action