

Oluwatobi Melvyn Mayungbo

Machine Learning Engineer

melvyn.t@gmail.com | [LinkedIn: linkedin.com/in/oluwatobi-mayungbo-3a567026b](https://www.linkedin.com/in/oluwatobi-mayungbo-3a567026b) | [GitHub: github.com/MelvTheGoat](https://github.com/MelvTheGoat)

SUMMARY

Machine learning engineer who builds and deploys production ML. Built a demand forecasting platform with backtesting, drift monitoring and cost-optimal forecasts, a fraud model served at 5.1 ms p99, and an uplift model for targeting, all deployed with FastAPI and Docker.

SKILLS

Payments and Risk: Demand Forecasting, Transaction Fraud Detection, Credit Risk Decisioning, Payment Reconciliation, Double-Entry Ledgers and Idempotency, Regulatory Compliance (CBN/NDPA), Web3 Fraud Risk Scoring

Machine Learning: LightGBM, PyTorch, XGBoost, Feature Engineering, Scikit-learn, Deep Learning, Sequence Models (GRU, TCN, Transformers), Model Calibration (Isotonic Regression, Temperature Scaling), SHAP

Engineering and MLOps: Python, FastAPI, Docker, Prefect, MLflow, CI/CD (GitHub Actions), SQL, DuckDB, PostgreSQL, Redis, GCP Cloud Run, AWS ECS Fargate, Railway, dbt-core, Evidently, Git, Pytest, Streamlit, React

Statistics and Causal Inference: Experiment Design, Uplift Modeling, Statistical Inference, Power Analysis, CUPED, Causal Inference (S/T/X-Learners, Causal Forests), Qini and AUUC, Time-Series Backtesting, Logistic Regression, WOE Scorecards

Data Analysis: Pandas, NumPy, SQL, Anomaly Detection

AI and LLM Systems: Retrieval-Augmented Generation (RAG), Hybrid Search (BM25, Dense Embeddings, RRF), LLM Evaluation and Guardrails, Prompt Engineering, Citation Verification, Model Context Protocol (MCP), Hugging Face, ONNX Runtime

PROJECTS

Self-Operating Demand Forecasting Platform | [GitHub](#)

Tech: Python, Prefect, dbt-core, DuckDB, LightGBM, MLflow, Evidently, FastAPI, Docker, GitHub Actions

- Built an end-to-end hourly demand forecasting system for NYC taxi zones, producing 24-hour quantile forecasts through an orchestrated Prefect pipeline spanning ingestion, dbt/DuckDB transformations, backtesting, drift monitoring and conditional retraining.
- Modelled asymmetric business cost by penalising unmet demand at 3× idle supply and shipping the cost-optimal quantile rather than the median, cutting expected cost by 25% on an identical model and beating a seasonal-naive baseline by 32%.
- Validated with rolling-origin backtesting across 6 expanding-window folds against seasonal-naive and historical-mean baselines, reporting MASE and pinball loss, with automated leakage tests that fail the build if any feature reads past the forecast origin.

Sequence-Based Transaction Fraud Detection System | [GitHub](#)

Tech: Python, PyTorch, LightGBM, ONNX Runtime, FastAPI, Docker, SHAP

- Benchmarked all architectures against a strong LightGBM baseline on engineered velocity and deviation features, reporting PR-AUC and precision at fixed review capacity (top 0.1%, 0.5%, 1%) across 3 seeds, with a per-attack-type breakdown covering card testing, account takeover and merchant compromise; sequence models cut expected cost per transaction 27% under an explicit cost model.
- Deployed real-time scoring by exporting to ONNX Runtime (2.9× faster than eager PyTorch at batch size 1) and serving via FastAPI, reaching 5.1 ms p99 end-to-end latency against a 50 ms budget, with append-only per-decision audit logging and attention-based attribution of the prior transactions driving each flag.

Uplift Modelling for Campaign Targeting | [GitHub](#)

Tech: Python, scikit-learn, LightGBM, EconML, pytest, GitHub Actions

- Ran a causal inference study on the Hillstrom randomised e-mail trial (64,000 customers) to decide which customers to contact under a fixed budget, delivered as a two-page decision memo for a marketing director.
- Showed the campaign works (1.25% conversion against 0.57%) but that a non-randomised list overstates the effect by 7.5×, and that a response model mails five times as many customers the campaign harms as an uplift model does.

- Added a power and MDE calculator and a CUPED variance-reduction demo; 135 tests on synthetic data with CI on Python 3.10 to 3.12.

Telco Customer Churn, End to End on AWS | [GitHub](#)

Tech: Python, XGBoost, MLflow, FastAPI, Gradio, Docker, GitHub Actions, AWS ECS Fargate, Application Load Balancer, CloudWatch

- Served it through a FastAPI /predict endpoint with a Gradio UI, containerised with Docker and deployed on AWS ECS Fargate behind an Application Load Balancer, with CloudWatch logs and GitHub Actions CI/CD.

EXPERIENCE

Machine Learning Instructor, SQI College of ICT March 2026 to Present

- Teach machine learning and deep learning to student cohorts, covering statistical foundations, classical ML algorithms, neural architectures, sequence models and production evaluation metrics.
- Instruct students on data analysis, feature engineering and model building with PyTorch, Scikit-Learn, Pandas and NumPy.

Internal Auditor, NISER May 2024 to July 2024

- Analysed financial and operational records to identify inconsistencies and irregularities, performing manual anomaly detection across complex transactional datasets.

EDUCATION

B.Sc. Statistics, University of Ibadan January 2021 to March 2025

Professional Diploma in Artificial Intelligence, SQI College of ICT July 2025 to Present

CERTIFICATIONS

- Google Data Analytics Professional Certificate
- Google Advanced Data Analytics Professional Certificate
- Claude Code In Action