

Oluwatobi Melvyn Mayungbo

Machine Learning and AI Engineer

melvyn.t@gmail.com | [LinkedIn: linkedin.com/in/oluwatobi-mayungbo-3a567026b](https://www.linkedin.com/in/oluwatobi-mayungbo-3a567026b) | [GitHub: github.com/MelvTheGoat](https://github.com/MelvTheGoat)

SUMMARY

Machine learning and AI engineer who builds classical ML and agent tools together. Built fraud detection, demand forecasting and cost-based decision models, and an MCP server of tools for AI assistants. Ships with testing, CI and monitoring, including drift detection and automated retraining.

SKILLS

Machine Learning: LightGBM, PyTorch, Feature Engineering, Scikit-learn, XGBoost, Deep Learning, Sequence Models (GRU, TCN, Transformers), Model Calibration (Isotonic Regression, Temperature Scaling), SHAP

Payments and Risk: Demand Forecasting, Transaction Fraud Detection, Credit Risk Decisioning, Payment Reconciliation, Double-Entry Ledgers and Idempotency, Regulatory Compliance (CBN/NDPA), Web3 Fraud Risk Scoring

AI and LLM Systems: Model Context Protocol (MCP), LLM Evaluation and Guardrails, Retrieval-Augmented Generation (RAG), Hybrid Search (BM25, Dense Embeddings, RRF), Prompt Engineering, Citation Verification, Hugging Face, ONNX Runtime

Engineering and MLOps: Python, MLflow, Evidently, CI/CD (GitHub Actions), SQL, DuckDB, PostgreSQL, Redis, FastAPI, Docker, GCP Cloud Run, AWS ECS Fargate, Railway, Prefect, dbt-core, Git, Pytest, Streamlit, React

Statistics and Causal Inference: Statistical Inference, Experiment Design, Power Analysis, CUPED, Causal Inference (S/T/X-Learners, Causal Forests), Uplift Modeling, Qini and AUUC, Time-Series Backtesting, Logistic Regression, WOE Scorecards

Data Analysis: Pandas, NumPy, SQL, Anomaly Detection

PROJECTS

Self-Operating Demand Forecasting Platform | [GitHub](#)

Tech: Python, Prefect, dbt-core, DuckDB, LightGBM, MLflow, Evidently, FastAPI, Docker, GitHub Actions

- Built an end-to-end hourly demand forecasting system for NYC taxi zones, producing 24-hour quantile forecasts through an orchestrated Prefect pipeline spanning ingestion, dbt/DuckDB transformations, backtesting, drift monitoring and conditional retraining.
- Modelled asymmetric business cost by penalising unmet demand at 3× idle supply and shipping the cost-optimal quantile rather than the median, cutting expected cost by 25% on an identical model and beating a seasonal-naive baseline by 32%.
- Validated with rolling-origin backtesting across 6 expanding-window folds against seasonal-naive and historical-mean baselines, reporting MASE and pinball loss, with automated leakage tests that fail the build if any feature reads past the forecast origin.

Sequence-Based Transaction Fraud Detection System | [GitHub](#)

Tech: Python, PyTorch, LightGBM, ONNX Runtime, FastAPI, Docker, SHAP

- Benchmarked all architectures against a strong LightGBM baseline on engineered velocity and deviation features, reporting PR-AUC and precision at fixed review capacity (top 0.1%, 0.5%, 1%) across 3 seeds, with a per-attack-type breakdown covering card testing, account takeover and merchant compromise; sequence models cut expected cost per transaction 27% under an explicit cost model.
- Deployed real-time scoring by exporting to ONNX Runtime (2.9× faster than eager PyTorch at batch size 1) and serving via FastAPI, reaching 5.1 ms p99 end-to-end latency against a 50 ms budget, with append-only per-decision audit logging and attention-based attribution of the prior transactions driving each flag.

web3-risk-mcp: Web3 Fraud Risk MCP Server | [GitHub](#)

Tech: Python, Model Context Protocol (MCP), Etherscan V2, GoPlus, DexScreener, Docker, GitHub Actions

- Built an MCP server that gives AI assistants (Claude Desktop, Cursor) six read-only tools to check a crypto wallet, token or smart contract for risk before anyone interacts with it, returning a 0 to 100 risk score with a reason for every point, across Ethereum, Base, Arbitrum One, Polygon PoS and BNB Chain.
- Built a resilient HTTP layer with caching, per-source rate limits and retries with exponential backoff and jitter; one failing source never breaks an investigation, and missing data lowers confidence rather than the score.

Uplift Modelling for Campaign Targeting | [GitHub](#)

Tech: Python, scikit-learn, LightGBM, EconML, pytest, GitHub Actions

- Ran a causal inference study on the Hillstrom randomised e-mail trial (64,000 customers) to decide which customers to contact under a fixed budget, delivered as a two-page decision memo for a marketing director.

EXPERIENCE

Machine Learning Instructor, SQI College of ICT

March 2026 to Present

- Teach machine learning and deep learning to student cohorts, covering statistical foundations, classical ML algorithms, neural architectures, sequence models and production evaluation metrics.
- Instruct students on data analysis, feature engineering and model building with PyTorch, Scikit-Learn, Pandas and NumPy.

Internal Auditor, NISER

May 2024 to July 2024

- Analysed financial and operational records to identify inconsistencies and irregularities, performing manual anomaly detection across complex transactional datasets.

EDUCATION

B.Sc. Statistics, University of Ibadan

January 2021 to March 2025

Professional Diploma in Artificial Intelligence, SQI College of ICT

July 2025 to Present

CERTIFICATIONS

- Google Data Analytics Professional Certificate
- Google Advanced Data Analytics Professional Certificate
- Claude Code In Action