

Oluwatobi Melvyn Mayungbo

AI Engineer, LLM and Retrieval Systems

mlvyn.t@gmail.com | [LinkedIn: linkedin.com/in/oluwatobi-mayungbo-3a567026b](https://www.linkedin.com/in/oluwatobi-mayungbo-3a567026b) | [GitHub: github.com/MelvTheGoat](https://github.com/MelvTheGoat)

SUMMARY

ML engineer focused on search and ranking quality. Built hybrid retrieval with BM25 and dense embeddings fused by reciprocal rank fusion at 0.96 recall@5 and 0.842 MRR, structure-aware chunking, and ranking-style models judged on precision at fixed capacity. Ships models in production with FastAPI.

SKILLS

AI and LLM Systems: Hybrid Search (BM25, Dense Embeddings, RRF), Retrieval-Augmented Generation (RAG), LLM Evaluation and Guardrails, Hugging Face, Prompt Engineering, Citation Verification, Model Context Protocol (MCP), ONNX Runtime

Machine Learning: LightGBM, PyTorch, Feature Engineering, Scikit-learn, XGBoost, Deep Learning, Sequence Models (GRU, TCN, Transformers), Model Calibration (Isotonic Regression, Temperature Scaling), SHAP

Engineering and MLOps: Python, FastAPI, Docker, SQL, DuckDB, PostgreSQL, Redis, GCP Cloud Run, AWS ECS Fargate, Railway, Prefect, dbt-core, MLflow, Evidently, CI/CD (GitHub Actions), Git, Pytest, Streamlit, React

Statistics and Causal Inference: Statistical Inference, Experiment Design, Power Analysis, CUPED, Causal Inference (S/T/X-Learners, Causal Forests), Uplift Modeling, Qini and AUUC, Time-Series Backtesting, Logistic Regression, WOE Scorecards

Data Analysis: Pandas, NumPy, SQL, Anomaly Detection

PROJECTS

Fintech Compliance RAG Assistant | [GitHub](#) | [Live App](#)

Tech: Python, Streamlit, BM25, ONNX Runtime, Reciprocal Rank Fusion, FastAPI, Docker, GCP Cloud Run

- Built and deployed a regulatory question-answering system over central bank (CBN) circulars and the NDPA, combining sparse (BM25) and dense (ONNX-served MiniLM) retrieval fused via reciprocal rank fusion, reaching 0.96 recall@5 and 0.842 MRR on a 50-question labelled golden set with zero retrieval misses in the top 10.
- Engineered section-aware chunking that never merges across a document's own structural boundaries, keeping every citation attributable to exactly one clause, and validated it against a structure-blind fixed-window baseline rather than asserting it was better.
- Ran retrieval and generation evaluation as separate metrics to localise failure by component; traced a 0.40 refusal-correctness result to a lexical-overlap limitation in the offline stub judge and flagged re-validation against a live provider as the next step instead of overstating readiness.

Sequence-Based Transaction Fraud Detection System | [GitHub](#)

Tech: Python, PyTorch, LightGBM, ONNX Runtime, FastAPI, Docker, SHAP

- Benchmarked all architectures against a strong LightGBM baseline on engineered velocity and deviation features, reporting PR-AUC and precision at fixed review capacity (top 0.1%, 0.5%, 1%) across 3 seeds, with a per-attack-type breakdown covering card testing, account takeover and merchant compromise; sequence models cut expected cost per transaction 27% under an explicit cost model.
- Deployed real-time scoring by exporting to ONNX Runtime (2.9× faster than eager PyTorch at batch size 1) and serving via FastAPI, reaching 5.1 ms p99 end-to-end latency against a 50 ms budget, with append-only per-decision audit logging and attention-based attribution of the prior transactions driving each flag.

Premier League Match Predictor | [GitHub](#)

Tech: Python, LightGBM, scikit-learn, pandas, SQLite, Flask, Vercel

- Blended a LightGBM classifier over a 204-column feature table with a multinomial logistic regression, trained on log loss so draws keep realistic probability.
- Backtested gameweek by gameweek over 1,050 matches: 52.1% outcome accuracy against 43.2% for always picking the home side, and log loss 0.9948 against 1.0061 for an Elo-only baseline, with well calibrated probabilities.

EXPERIENCE

Machine Learning Instructor, SQI College of ICT

March 2026 to Present

- Teach machine learning and deep learning to student cohorts, covering statistical foundations, classical ML algorithms, neural architectures, sequence models and production evaluation metrics.
 - Instruct students on data analysis, feature engineering and model building with PyTorch, Scikit-Learn, Pandas and NumPy.
- Internal Auditor, NISER** May 2024 to July 2024
- Analysed financial and operational records to identify inconsistencies and irregularities, performing manual anomaly detection across complex transactional datasets.

EDUCATION

- B.Sc. Statistics, University of Ibadan** January 2021 to March 2025
- Professional Diploma in Artificial Intelligence, SQI College of ICT** July 2025 to Present

CERTIFICATIONS

- Google Data Analytics Professional Certificate
- Google Advanced Data Analytics Professional Certificate
- Claude Code In Action