

Oluwatobi Melvyn Mayungbo

AI Engineer, LLM and Retrieval Systems

mlvyn.t@gmail.com | [LinkedIn: linkedin.com/in/oluwatobi-mayungbo-3a567026b](https://www.linkedin.com/in/oluwatobi-mayungbo-3a567026b) | [GitHub: github.com/MelvTheGoat](https://github.com/MelvTheGoat)

SUMMARY

AI engineer with a B.Sc. in Statistics who builds evaluations and analyses how AI systems fail. Built an evaluation harness that scores retrieval and generation separately on a labelled golden set, a record and replay harness for an MCP tool server, and benchmarks reported across seeds against strong baselines.

SKILLS

AI and LLM Systems: LLM Evaluation and Guardrails, Retrieval-Augmented Generation (RAG), Model Context Protocol (MCP), Prompt Engineering, Hybrid Search (BM25, Dense Embeddings, RRF), Citation Verification, Hugging Face, ONNX Runtime

Statistics and Causal Inference: Statistical Inference, Experiment Design, Power Analysis, CUPED, Causal Inference (S/T/X-Learners, Causal Forests), Uplift Modeling, Qini and AUUC, Time-Series Backtesting, Logistic Regression, WOE Scorecards

Engineering and MLOps: Python, SQL, Pytest, CI/CD (GitHub Actions), DuckDB, PostgreSQL, Redis, FastAPI, Docker, GCP Cloud Run, AWS ECS Fargate, Railway, Prefect, dbt-core, MLflow, Evidently, Git, Streamlit, React

Machine Learning: PyTorch, Scikit-learn, LightGBM, XGBoost, Deep Learning, Sequence Models (GRU, TCN, Transformers), Model Calibration (Isotonic Regression, Temperature Scaling), SHAP, Feature Engineering

Payments and Risk: Transaction Fraud Detection, Credit Risk Decisioning, Payment Reconciliation, Double-Entry Ledgers and Idempotency, Regulatory Compliance (CBN/NDPA), Demand Forecasting, Web3 Fraud Risk Scoring

PROJECTS

Fintech Compliance RAG Assistant | [GitHub](#) | [Live App](#)

Tech: Python, Streamlit, BM25, ONNX Runtime, Reciprocal Rank Fusion, FastAPI, Docker, GCP Cloud Run

- Ran retrieval and generation evaluation as separate metrics to localise failure by component; traced a 0.40 refusal-correctness result to a lexical-overlap limitation in the offline stub judge and flagged re-validation against a live provider as the next step instead of overstating readiness.
- Designed generation to fail safely under uncertainty: a sentinel-token refusal mechanism enforced across the prompt, output guardrail and evaluation harness, plus post-generation citation verification that strips any marker the model cited but never retrieved.
- Built and deployed a regulatory question-answering system over central bank (CBN) circulars and the NDPA, combining sparse (BM25) and dense (ONNX-served MiniLM) retrieval fused via reciprocal rank fusion, reaching 0.96 recall@5 and 0.842 MRR on a 50-question labelled golden set with zero retrieval misses in the top 10.

web3-risk-mcp: Web3 Fraud Risk MCP Server | [GitHub](#)

Tech: Python, Model Context Protocol (MCP), Etherscan V2, GoPlus, DexScreener, Docker, GitHub Actions

- Assembled an evaluation set of 34 hand-checked addresses (12 risky, 22 safe) with a record and replay harness reporting ROC AUC, precision and recall.
- Built an MCP server that gives AI assistants (Claude Desktop, Cursor) six read-only tools to check a crypto wallet, token or smart contract for risk before anyone interacts with it, returning a 0 to 100 risk score with a reason for every point, across Ethereum, Base, Arbitrum One, Polygon PoS and BNB Chain.

Sequence-Based Transaction Fraud Detection System | [GitHub](#)

Tech: Python, PyTorch, LightGBM, ONNX Runtime, FastAPI, Docker, SHAP

- Benchmarked all architectures against a strong LightGBM baseline on engineered velocity and deviation features, reporting PR-AUC and precision at fixed review capacity (top 0.1%, 0.5%, 1%) across 3 seeds, with a per-attack-type breakdown covering card testing, account takeover and merchant compromise; sequence models cut expected cost per transaction 27% under an explicit cost model.
- Handled extreme class imbalance with focal loss and weighted sampling and isolated their failure modes: focal loss degraded probability calibration (ECE 0.0063) while weighted sampling cost 0.042 PR-AUC. Corrected calibration with temperature scaling and isotonic regression, reducing ECE to 0.0005.

EXPERIENCE

Machine Learning Instructor, SQI College of ICT

March 2026 to Present

- Teach machine learning and deep learning to student cohorts, covering statistical foundations, classical ML algorithms, neural architectures, sequence models and production evaluation metrics.
- Instruct students on data analysis, feature engineering and model building with PyTorch, Scikit-Learn, Pandas and NumPy.

Internal Auditor, NISER

May 2024 to July 2024

- Analysed financial and operational records to identify inconsistencies and irregularities, performing manual anomaly detection across complex transactional datasets.

EDUCATION

B.Sc. Statistics, University of Ibadan

January 2021 to March 2025

Professional Diploma in Artificial Intelligence, SQI College of ICT

July 2025 to Present

CERTIFICATIONS

- Google Data Analytics Professional Certificate
- Google Advanced Data Analytics Professional Certificate
- Claude Code In Action