

Oluwatobi Melvyn Mayungbo

Data Scientist, Machine Learning and Statistics

mlvyn.t@gmail.com | [LinkedIn: linkedin.com/in/oluwatobi-mayungbo-3a567026b](https://www.linkedin.com/in/oluwatobi-mayungbo-3a567026b) | [GitHub: github.com/MelvTheGoat](https://github.com/MelvTheGoat)

SUMMARY

Data scientist with a B.Sc. in Statistics focused on experiments and root-cause analysis. Evaluated targeting on a randomised campaign trial with 64,000 customers, built power, MDE and CUPED tools, and forecasts demand with rolling backtests. Explains results as decisions for product and business partners.

SKILLS

Statistics and Causal Inference: Experiment Design, Power Analysis, CUPED, Causal Inference (S/T/X-Learners, Causal Forests), Statistical Inference, Uplift Modeling, Qini and AUUC, Time-Series Backtesting, Logistic Regression, WOE Scorecards

Engineering and MLOps: SQL, Python, DuckDB, PostgreSQL, Redis, FastAPI, Docker, GCP Cloud Run, AWS ECS Fargate, Railway, Prefect, dbt-core, MLflow, Evidently, CI/CD (GitHub Actions), Git, Pytest, Streamlit, React

Data Analysis: Pandas, NumPy, SQL, Anomaly Detection

Machine Learning: PyTorch, Scikit-learn, LightGBM, XGBoost, Deep Learning, Sequence Models (GRU, TCN, Transformers), Model Calibration (Isotonic Regression, Temperature Scaling), SHAP, Feature Engineering

AI and LLM Systems: Retrieval-Augmented Generation (RAG), Hybrid Search (BM25, Dense Embeddings, RRF), LLM Evaluation and Guardrails, Prompt Engineering, Citation Verification, Model Context Protocol (MCP), Hugging Face, ONNX Runtime

Payments and Risk: Transaction Fraud Detection, Credit Risk Decisioning, Payment Reconciliation, Double-Entry Ledgers and Idempotency, Regulatory Compliance (CBN/NDPA), Demand Forecasting, Web3 Fraud Risk Scoring

PROJECTS

Uplift Modelling for Campaign Targeting | [GitHub](#)

Tech: Python, scikit-learn, LightGBM, EconML, pytest, GitHub Actions

- Ran a causal inference study on the Hillstrom randomised e-mail trial (64,000 customers) to decide which customers to contact under a fixed budget, delivered as a two-page decision memo for a marketing director.
- Showed the campaign works (1.25% conversion against 0.57%) but that a non-randomised list overstates the effect by 7.5×, and that a response model mails five times as many customers the campaign harms as an uplift model does.
- Evaluated with Qini curves, AUUC, decile tables and bootstrap intervals; uplift targeting was worth about \$1,390 against about \$3,170 for raising the budget, and a failed placebo test on the primary campaign was reported openly.
- Added a power and MDE calculator and a CUPED variance-reduction demo; 135 tests on synthetic data with CI on Python 3.10 to 3.12.

Self-Operating Demand Forecasting Platform | [GitHub](#)

Tech: Python, Prefect, dbt-core, DuckDB, LightGBM, MLflow, Evidently, FastAPI, Docker, GitHub Actions

- Built an end-to-end hourly demand forecasting system for NYC taxi zones, producing 24-hour quantile forecasts through an orchestrated Prefect pipeline spanning ingestion, dbt/DuckDB transformations, backtesting, drift monitoring and conditional retraining.
- Validated with rolling-origin backtesting across 6 expanding-window folds against seasonal-naïve and historical-mean baselines, reporting MASE and pinball loss, with automated leakage tests that fail the build if any feature reads past the forecast origin.

Telco Customer Churn, End to End on AWS | [GitHub](#)

Tech: Python, XGBoost, MLflow, FastAPI, Gradio, Docker, GitHub Actions, AWS ECS Fargate, Application Load Balancer, CloudWatch

- Built a churn prediction model with feature engineering and an XGBoost classifier, with experiments tracked in MLflow.

Premier League Match Predictor | [GitHub](#)

Tech: Python, LightGBM, scikit-learn, pandas, SQLite, Flask, Vercel

- Backtested gameweek by gameweek over 1,050 matches: 52.1% outcome accuracy against 43.2% for always picking the home side, and log loss 0.9948 against 1.0061 for an Elo-only baseline, with well calibrated probabilities.

EXPERIENCE

Machine Learning Instructor, SQI College of ICT

March 2026 to Present

- Teach machine learning and deep learning to student cohorts, covering statistical foundations, classical ML algorithms, neural architectures, sequence models and production evaluation metrics.
- Instruct students on data analysis, feature engineering and model building with PyTorch, Scikit-Learn, Pandas and NumPy.

Internal Auditor, NISER

May 2024 to July 2024

- Analysed financial and operational records to identify inconsistencies and irregularities, performing manual anomaly detection across complex transactional datasets.

EDUCATION

B.Sc. Statistics, University of Ibadan

January 2021 to March 2025

Professional Diploma in Artificial Intelligence, SQI College of ICT

July 2025 to Present

CERTIFICATIONS

- Google Data Analytics Professional Certificate
- Google Advanced Data Analytics Professional Certificate
- Claude Code In Action