

Oluwatobi Melvyn Mayungbo

Data Scientist, Credit and Fraud Risk

mlvyn.t@gmail.com | [LinkedIn: linkedin.com/in/oluwatobi-mayungbo-3a567026b](https://www.linkedin.com/in/oluwatobi-mayungbo-3a567026b) | [GitHub: github.com/MelvTheGoat](https://github.com/MelvTheGoat)

SUMMARY

Credit risk data scientist with a B.Sc. in Statistics. Built a credit decisioning system with calibrated default probabilities, cost-optimal cutoffs, reason codes and a fairness audit, and compared a WOE-binned logistic scorecard with gradient boosting. Runs policy trade-offs as explicit cost models.

SKILLS

Payments and Risk: Credit Risk Decisioning, Transaction Fraud Detection, Payment Reconciliation, Double-Entry Ledgers and Idempotency, Regulatory Compliance (CBN/NDPA), Demand Forecasting, Web3 Fraud Risk Scoring

Statistics and Causal Inference: WOE Scorecards, Logistic Regression, Experiment Design, Statistical Inference, Power Analysis, CUPED, Causal Inference (S/T/X-Learners, Causal Forests), Uplift Modeling, Qini and AUUC, Time-Series Backtesting

Machine Learning: LightGBM, Model Calibration (Isotonic Regression, Temperature Scaling), SHAP, PyTorch, Scikit-learn, XGBoost, Deep Learning, Sequence Models (GRU, TCN, Transformers), Feature Engineering

Data Analysis: SQL, Pandas, NumPy, Anomaly Detection

Engineering and MLOps: Python, SQL, DuckDB, PostgreSQL, Redis, FastAPI, Docker, GCP Cloud Run, AWS ECS Fargate, Railway, Prefect, dbt-core, MLflow, Evidently, CI/CD (GitHub Actions), Git, Pytest, Streamlit, React

AI and LLM Systems: Retrieval-Augmented Generation (RAG), Hybrid Search (BM25, Dense Embeddings, RRF), LLM Evaluation and Guardrails, Prompt Engineering, Citation Verification, Model Context Protocol (MCP), Hugging Face, ONNX Runtime

PROJECTS

Credit Risk Decisioning and Fairness Audit System | [GitHub](#) | [Live App](#)

Tech: Python, LightGBM, scikit-learn, SHAP, WOE Scorecards, FastAPI, Docker

- Developed a full credit decisioning system producing calibrated default probabilities, cost-optimal approve and decline cutoffs, adverse action reason codes and an append-only audit trail that lets any individual decision be reconstructed months later.
- Benchmarked gradient boosting against a traditional WOE-binned logistic scorecard, prioritising calibration quality (Brier score, reliability diagrams, expected calibration error) over ranking metrics, because expected loss needs true probabilities rather than rankings.
- Audited fairness across sex, age, education and marital status using demographic parity and equal opportunity differences plus within-group calibration, and presented the accuracy and fairness tradeoff curve as a policy decision rather than silently optimising it away.

Uplift Modelling for Campaign Targeting | [GitHub](#)

Tech: Python, scikit-learn, LightGBM, EconML, pytest, GitHub Actions

- Showed the campaign works (1.25% conversion against 0.57%) but that a non-randomised list overstates the effect by 7.5×, and that a response model mails five times as many customers the campaign harms as an uplift model does.
- Evaluated with Qini curves, AUUC, decile tables and bootstrap intervals; uplift targeting was worth about \$1,390 against about \$3,170 for raising the budget, and a failed placebo test on the primary campaign was reported openly.

Self-Operating Demand Forecasting Platform | [GitHub](#)

Tech: Python, Prefect, dbt-core, DuckDB, LightGBM, MLflow, Evidently, FastAPI, Docker, GitHub Actions

- Modelled asymmetric business cost by penalising unmet demand at 3× idle supply and shipping the cost-optimal quantile rather than the median, cutting expected cost by 25% on an identical model and beating a seasonal-naive baseline by 32%.

EXPERIENCE

Machine Learning Instructor, SQI College of ICT

March 2026 to Present

- Teach machine learning and deep learning to student cohorts, covering statistical foundations, classical ML algorithms, neural architectures, sequence models and production evaluation metrics.

- Instruct students on data analysis, feature engineering and model building with PyTorch, Scikit-Learn, Pandas and NumPy.
- Internal Auditor, NISER**

May 2024 to July 2024
- Analysed financial and operational records to identify inconsistencies and irregularities, performing manual anomaly detection across complex transactional datasets.

EDUCATION

- B.Sc. Statistics, University of Ibadan**

January 2021 to March 2025
- Professional Diploma in Artificial Intelligence, SQI College of ICT**

July 2025 to Present

CERTIFICATIONS

- Google Data Analytics Professional Certificate
- Google Advanced Data Analytics Professional Certificate
- Claude Code In Action