

Oluwatobi Melvyn Mayungbo

Data Analyst, Fraud and Risk

mlvyn.t@gmail.com | [LinkedIn: linkedin.com/in/oluwatobi-mayungbo-3a567026b](https://www.linkedin.com/in/oluwatobi-mayungbo-3a567026b) | [GitHub: github.com/MelvTheGoat](https://github.com/MelvTheGoat)

SUMMARY

Data analyst and machine learning engineer with a B.Sc. in Statistics, focused on fraud and payments data. Built a sequence-based fraud detection system analysed by attack type, and a reconciliation system that finds duplicate settlements and payout discrepancies. Earlier did manual anomaly detection on transactional records as an internal auditor. Works mainly in SQL and Python.

SKILLS

Payments and Risk: Transaction Fraud Detection, Payment Reconciliation, Credit Risk Decisioning, Double-Entry Ledgers and Idempotency, Regulatory Compliance (CBN/NDPA), Demand Forecasting, Web3 Fraud Risk Scoring

Data Analysis: SQL, Anomaly Detection, Pandas, NumPy

Machine Learning: PyTorch, Scikit-learn, LightGBM, XGBoost, Deep Learning, Sequence Models (GRU, TCN, Transformers), Model Calibration (Isotonic Regression, Temperature Scaling), SHAP, Feature Engineering

Statistics and Causal Inference: Statistical Inference, Experiment Design, Power Analysis, CUPED, Causal Inference (S/T/X-Learners, Causal Forests), Uplift Modeling, Qini and AUUC, Time-Series Backtesting, Logistic Regression, WOE Scorecards

Engineering and MLOps: Python, SQL, DuckDB, PostgreSQL, Redis, FastAPI, Docker, GCP Cloud Run, AWS ECS Fargate, Railway, Prefect, dbt-core, MLflow, Evidently, CI/CD (GitHub Actions), Git, Pytest, Streamlit, React

AI and LLM Systems: Retrieval-Augmented Generation (RAG), Hybrid Search (BM25, Dense Embeddings, RRF), LLM Evaluation and Guardrails, Prompt Engineering, Citation Verification, Model Context Protocol (MCP), Hugging Face, ONNX Runtime

PROJECTS

Sequence-Based Transaction Fraud Detection System | [GitHub](#)

Tech: Python, PyTorch, LightGBM, ONNX Runtime, FastAPI, Docker, SHAP

- Benchmarked all architectures against a strong LightGBM baseline on engineered velocity and deviation features, reporting PR-AUC and precision at fixed review capacity (top 0.1%, 0.5%, 1%) across 3 seeds, with a per-attack-type breakdown covering card testing, account takeover and merchant compromise; sequence models cut expected cost per transaction 27% under an explicit cost model.
- Handled extreme class imbalance with focal loss and weighted sampling and isolated their failure modes: focal loss degraded probability calibration (ECE 0.0063) while weighted sampling cost 0.042 PR-AUC. Corrected calibration with temperature scaling and isotonic regression, reducing ECE to 0.0005.
- Built and compared three deep sequence architectures (GRU, Temporal Convolutional Network and a causally masked Transformer encoder) over customer transaction histories, with learned embeddings for high-cardinality merchant (~2.9k vocabulary), MCC and device-type categoricals; all models under 115k parameters and CPU-trainable in under 6 minutes.

Reckon: Payment Reconciliation and Review System | [GitHub](#) | [Live App](#)

Tech: Python, PyTorch, DuckDB, FastAPI, React, Docker, Railway, PostgreSQL

- Built and deployed a payment reconciliation system processing multi-source transaction feeds to detect uncollected revenue, duplicate settlements and payout discrepancies, combining deterministic matching rules with ML-assisted discrepancy resolution.
- Engineered a multi-stage reconciliation pipeline using DuckDB and FastAPI, with temporal join logic and tolerance thresholds to handle asynchronous processor webhooks and bank settlement files.
- Designed a human-in-the-loop review interface in React showing match confidence scores, audit trails and suggested resolution actions to streamline exception handling.

Credit Risk Decisioning and Fairness Audit System | [GitHub](#) | [Live App](#)

Tech: Python, LightGBM, scikit-learn, SHAP, WOE Scorecards, FastAPI, Docker

- Developed a full credit decisioning system producing calibrated default probabilities, cost-optimal approve and decline cutoffs, adverse action reason codes and an append-only audit trail that lets any individual decision be reconstructed months later.
- Benchmarked gradient boosting against a traditional WOE-binned logistic scorecard, prioritising calibration quality (Brier score, reliability diagrams, expected calibration error) over ranking metrics, because expected loss needs true probabilities rather than rankings.

Self-Operating Demand Forecasting Platform | [GitHub](#)

Tech: Python, Prefect, dbt-core, DuckDB, LightGBM, MLflow, Evidently, FastAPI, Docker, GitHub Actions

- Built an end-to-end hourly demand forecasting system for NYC taxi zones, producing 24-hour quantile forecasts through an orchestrated Prefect pipeline spanning ingestion, dbt/DuckDB transformations, backtesting, drift monitoring and conditional retraining.

EXPERIENCE

Machine Learning Instructor, SQI College of ICT

March 2026 to Present

- Teach machine learning and deep learning to student cohorts, covering statistical foundations, classical ML algorithms, neural architectures, sequence models and production evaluation metrics.
- Instruct students on data analysis, feature engineering and model building with PyTorch, Scikit-Learn, Pandas and NumPy.

Internal Auditor, NISER

May 2024 to July 2024

- Analysed financial and operational records to identify inconsistencies and irregularities, performing manual anomaly detection across complex transactional datasets.

EDUCATION

B.Sc. Statistics, University of Ibadan

January 2021 to March 2025

Professional Diploma in Artificial Intelligence, SQI College of ICT

July 2025 to Present

CERTIFICATIONS

- Google Data Analytics Professional Certificate
- Google Advanced Data Analytics Professional Certificate
- Claude Code In Action