

Oluwatobi Melvyn Mayungbo

AI Engineer, LLM and Retrieval Systems

melvyn.t@gmail.com | [LinkedIn: linkedin.com/in/oluwatobi-mayungbo-3a567026b](https://www.linkedin.com/in/oluwatobi-mayungbo-3a567026b) | [GitHub: github.com/MelvTheGoat](https://github.com/MelvTheGoat)

SUMMARY

AI engineer who turns LLM research ideas into measured prototypes. Built a retrieval assistant with hybrid search, guardrails and a labelled benchmark, an MCP server for AI assistants with a replay evaluation harness, and ML models benchmarked against strong baselines. Writes clean, tested Python.

SKILLS

AI and LLM Systems: Retrieval-Augmented Generation (RAG), Hybrid Search (BM25, Dense Embeddings, RRF), LLM Evaluation and Guardrails, Prompt Engineering, Model Context Protocol (MCP), Hugging Face, Citation Verification, ONNX Runtime

Engineering and MLOps: Python, Pytest, Docker, SQL, DuckDB, PostgreSQL, Redis, FastAPI, GCP Cloud Run, AWS ECS Fargate, Railway, Prefect, dbt-core, MLflow, Evidently, CI/CD (GitHub Actions), Git, Streamlit, React

Machine Learning: PyTorch, Scikit-learn, LightGBM, XGBoost, Deep Learning, Sequence Models (GRU, TCN, Transformers), Model Calibration (Isotonic Regression, Temperature Scaling), SHAP, Feature Engineering

Statistics and Causal Inference: Statistical Inference, Experiment Design, Power Analysis, CUPED, Causal Inference (S/T/X-Learners, Causal Forests), Uplift Modeling, Qini and AUUC, Time-Series Backtesting, Logistic Regression, WOE Scorecards

Data Analysis: Pandas, NumPy, SQL, Anomaly Detection

PROJECTS

Fintech Compliance RAG Assistant | [GitHub](#) | [Live App](#)

Tech: Python, Streamlit, BM25, ONNX Runtime, Reciprocal Rank Fusion, FastAPI, Docker, GCP Cloud Run

- Built and deployed a regulatory question-answering system over central bank (CBN) circulars and the NDPA, combining sparse (BM25) and dense (ONNX-served MiniLM) retrieval fused via reciprocal rank fusion, reaching 0.96 recall@5 and 0.842 MRR on a 50-question labelled golden set with zero retrieval misses in the top 10.
- Engineered section-aware chunking that never merges across a document's own structural boundaries, keeping every citation attributable to exactly one clause, and validated it against a structure-blind fixed-window baseline rather than asserting it was better.
- Ran retrieval and generation evaluation as separate metrics to localise failure by component; traced a 0.40 refusal-correctness result to a lexical-overlap limitation in the offline stub judge and flagged re-validation against a live provider as the next step instead of overstating readiness.
- Designed generation to fail safely under uncertainty: a sentinel-token refusal mechanism enforced across the prompt, output guardrail and evaluation harness, plus post-generation citation verification that strips any marker the model cited but never retrieved.

web3-risk-mcp: Web3 Fraud Risk MCP Server | [GitHub](#)

Tech: Python, Model Context Protocol (MCP), Etherscan V2, GoPlus, DexScreener, Docker, GitHub Actions

- Built an MCP server that gives AI assistants (Claude Desktop, Cursor) six read-only tools to check a crypto wallet, token or smart contract for risk before anyone interacts with it, returning a 0 to 100 risk score with a reason for every point, across Ethereum, Base, Arbitrum One, Polygon PoS and BNB Chain.
- Assembled an evaluation set of 34 hand-checked addresses (12 risky, 22 safe) with a record and replay harness reporting ROC AUC, precision and recall.

Sequence-Based Transaction Fraud Detection System | [GitHub](#)

Tech: Python, PyTorch, LightGBM, ONNX Runtime, FastAPI, Docker, SHAP

- Built and compared three deep sequence architectures (GRU, Temporal Convolutional Network and a causally masked Transformer encoder) over customer transaction histories, with learned embeddings for high-cardinality merchant (~2.9k vocabulary), MCC and device-type categoricals; all models under 115k parameters and CPU-trainable in under 6 minutes.
- Benchmarked all architectures against a strong LightGBM baseline on engineered velocity and deviation features, reporting PR-AUC and precision at fixed review capacity (top 0.1%, 0.5%, 1%) across 3 seeds, with a per-attack-type

breakdown covering card testing, account takeover and merchant compromise; sequence models cut expected cost per transaction 27% under an explicit cost model.

Self-Operating Demand Forecasting Platform | [GitHub](#)

Tech: Python, Prefect, dbt-core, DuckDB, LightGBM, MLflow, Evidently, FastAPI, Docker, GitHub Actions

- Validated with rolling-origin backtesting across 6 expanding-window folds against seasonal-naive and historical-mean baselines, reporting MASE and pinball loss, with automated leakage tests that fail the build if any feature reads past the forecast origin.

EXPERIENCE

Machine Learning Instructor, SQI College of ICT

March 2026 to Present

- Teach machine learning and deep learning to student cohorts, covering statistical foundations, classical ML algorithms, neural architectures, sequence models and production evaluation metrics.
- Instruct students on data analysis, feature engineering and model building with PyTorch, Scikit-Learn, Pandas and NumPy.

Internal Auditor, NISER

May 2024 to July 2024

- Analysed financial and operational records to identify inconsistencies and irregularities, performing manual anomaly detection across complex transactional datasets.

EDUCATION

B.Sc. Statistics, University of Ibadan

January 2021 to March 2025

Professional Diploma in Artificial Intelligence, SQI College of ICT

July 2025 to Present

CERTIFICATIONS

- Google Data Analytics Professional Certificate
- Google Advanced Data Analytics Professional Certificate
- Claude Code In Action