

Oluwatobi Melvyn Mayungbo

Machine Learning Engineer, Fraud and Risk

mlvyn.t@gmail.com | [LinkedIn: linkedin.com/in/oluwatobi-mayungbo-3a567026b](https://www.linkedin.com/in/oluwatobi-mayungbo-3a567026b) | [GitHub: github.com/MelvTheGoat](https://github.com/MelvTheGoat)

SUMMARY

Machine learning engineer with a B.Sc. in Statistics who builds deep learning fraud detection. Built GRU, TCN and Transformer models over customer transaction histories that beat a strong LightGBM baseline, cut expected cost per transaction 27% and serve in real time at 5.1 ms p99 through ONNX Runtime.

SKILLS

Machine Learning: PyTorch, Sequence Models (GRU, TCN, Transformers), LightGBM, Model Calibration (Isotonic Regression, Temperature Scaling), SHAP, Feature Engineering, Scikit-learn, XGBoost, Deep Learning

Payments and Risk: Transaction Fraud Detection, Web3 Fraud Risk Scoring, Credit Risk Decisioning, Payment Reconciliation, Double-Entry Ledgers and Idempotency, Regulatory Compliance (CBN/NDPA), Demand Forecasting

Engineering and MLOps: Python, SQL, FastAPI, Docker, DuckDB, PostgreSQL, Redis, GCP Cloud Run, AWS ECS Fargate, Railway, Prefect, dbt-core, MLflow, Evidently, CI/CD (GitHub Actions), Git, Pytest, Streamlit, React

AI and LLM Systems: ONNX Runtime, Retrieval-Augmented Generation (RAG), Hybrid Search (BM25, Dense Embeddings, RRF), LLM Evaluation and Guardrails, Prompt Engineering, Citation Verification, Model Context Protocol (MCP), Hugging Face

Statistics and Causal Inference: Statistical Inference, Experiment Design, Power Analysis, CUPED, Causal Inference (S/T/X-Learners, Causal Forests), Uplift Modeling, Qini and AUUC, Time-Series Backtesting, Logistic Regression, WOE Scorecards

Data Analysis: Pandas, NumPy, SQL, Anomaly Detection

PROJECTS

Sequence-Based Transaction Fraud Detection System | [GitHub](#)

Tech: Python, PyTorch, LightGBM, ONNX Runtime, FastAPI, Docker, SHAP

- Built and compared three deep sequence architectures (GRU, Temporal Convolutional Network and a causally masked Transformer encoder) over customer transaction histories, with learned embeddings for high-cardinality merchant (~2.9k vocabulary), MCC and device-type categoricals; all models under 115k parameters and CPU-trainable in under 6 minutes.
- Benchmarked all architectures against a strong LightGBM baseline on engineered velocity and deviation features, reporting PR-AUC and precision at fixed review capacity (top 0.1%, 0.5%, 1%) across 3 seeds, with a per-attack-type breakdown covering card testing, account takeover and merchant compromise; sequence models cut expected cost per transaction 27% under an explicit cost model.
- Handled extreme class imbalance with focal loss and weighted sampling and isolated their failure modes: focal loss degraded probability calibration (ECE 0.0063) while weighted sampling cost 0.042 PR-AUC. Corrected calibration with temperature scaling and isotonic regression, reducing ECE to 0.0005.
- Deployed real-time scoring by exporting to ONNX Runtime (2.9× faster than eager PyTorch at batch size 1) and serving via FastAPI, reaching 5.1 ms p99 end-to-end latency against a 50 ms budget, with append-only per-decision audit logging and attention-based attribution of the prior transactions driving each flag.

web3-risk-mcp: Web3 Fraud Risk MCP Server | [GitHub](#)

Tech: Python, Model Context Protocol (MCP), Etherscan V2, GoPlus, DexScreener, Docker, GitHub Actions

- Detects honeypot tokens, rug pull signs, hidden owner powers (mint, blacklist, pause), buy and sell tax, holder concentration, unlocked liquidity, and fund links to mixers, sanctioned wallets, exploiters and phishing addresses over 1 or 2 hops.
- Designed a rule-based, fully explainable scorer where findings carry stable IDs and decisive findings set a floor of 75 so trust signals can never hide them; each score carries a confidence level based on how many sources answered.

Credit Risk Decisioning and Fairness Audit System | [GitHub](#) | [Live App](#)

Tech: Python, LightGBM, scikit-learn, SHAP, WOE Scorecards, FastAPI, Docker

- Benchmarked gradient boosting against a traditional WOE-binned logistic scorecard, prioritising calibration quality (Brier score, reliability diagrams, expected calibration error) over ranking metrics, because expected loss needs true probabilities rather than rankings.

EXPERIENCE

- Machine Learning Instructor, SQI College of ICT**March 2026 to Present
- Teach machine learning and deep learning to student cohorts, covering statistical foundations, classical ML algorithms, neural architectures, sequence models and production evaluation metrics.
 - Instruct students on data analysis, feature engineering and model building with PyTorch, Scikit-Learn, Pandas and NumPy.
- Internal Auditor, NISER**May 2024 to July 2024
- Analysed financial and operational records to identify inconsistencies and irregularities, performing manual anomaly detection across complex transactional datasets.

EDUCATION

- B.Sc. Statistics, University of Ibadan**January 2021 to March 2025
- Professional Diploma in Artificial Intelligence, SQI College of ICT**July 2025 to Present

CERTIFICATIONS

- Google Data Analytics Professional Certificate
- Google Advanced Data Analytics Professional Certificate
- Claude Code In Action